

283 lines (197 loc) · 13.9 KB

PreviewCodeBlame

RawCopyDownloadEditDropdownMenu



caveman

why use many token when few do trick

stars74k

last commitlast friday

licenseinvalid

[Before/After](#) • [Install](#) • [What You Get](#) • [Benchmarks](#) • [Full install guide](#)

A [Claude Code](#) skill/plugin (also Codex, Gemini, Cursor, Windsurf, Cline, Copilot, 30+ more) that makes agent talk like caveman — cuts **~75% of output tokens**, keeps full technical accuracy. Brain still big. Mouth small.

Before / After

Normal Claude (69 tokens)	Caveman Claude (19 tokens) "New object ref each render. Inline object prop = new ref = re-render. Wrap in useMemo ."
---------------------------	---

"The reason your React component is re-rendering is likely because you're creating a new object reference on each render cycle. When you pass an inline object as a prop, React's shallow comparison sees it as a different object every time, which triggers a re-render. I'd recommend using `useMemo` to memoize the object."

Normal Claude

"Sure! I'd be happy to help you with that. The issue you're experiencing is most likely caused by your authentication middleware not properly validating the token expiry. Let me take a look and suggest a fix."

Caveman Claude

"Bug in auth middleware. Token expiry check use `<` not `<=` . Fix:"

Same fix. 75% less word. Brain still big.

TOKENS SAVED		75%
TECHNICAL ACCURACY		100%
SPEED INCREASE		~3x
VIBES		00G



Pick your level of grunt — `lite` (drop filler), `full` (default caveman), `ultra` (telegraphic), or `wenyan` (classical Chinese, even shorter). One command switch. Cost go down forever.

Speak your tongue. Caveman keep your language. You write Portuguese, caveman grunt Portuguese. Spanish, French, same. Compress the *style*, not the language. Code, command, error string stay exact.

"Novo ref de objeto cada render. Prop inline = novo ref = re-render. Envolve com `useMemo` ."

Like this trick? Now get whole agent — **caveman-code**

This skill shrink what agent say. [caveman-code](#) shrink **everything** — full terminal coding agent, caveman top to bottom. ~2× **fewer tokens than Codex** on identical tasks. 20+ providers · plan mode · autopilot goal loop · MIT.

```
npm install -g @juliusbrussee/caveman-code
```



[▶ Try caveman-code now →](#) — *why use many token when whole agent save*

Install

One line. Find every agent. Install for each.

```
# macOS / Linux / WSL / Git Bash
curl -fsSL https://raw.githubusercontent.com/JuliusBrussee/caveman/main/install.sh | bash

# Windows (PowerShell 5.1+)
irm https://raw.githubusercontent.com/JuliusBrussee/caveman/main/install.ps1 | iex
```

~30 seconds. Needs Node ≥18. Skip agent you no have. Safe to re-run.

Trigger: type `/caveman` or say "talk like caveman". Stop with "normal mode".

One agent only, manual command, or any of 30+ other agents → [INSTALL.md](#). Install break? Open agent, say "Read *CLAUDE.md* and *INSTALL.md*, install caveman for me." Agent fix own brain.

What You Get

Skill	What
<code>/caveman</code> <code>[lite full ultra wenyan]</code>	Compress every reply. Levels stick until session end.
<code>/caveman-commit</code>	Conventional Commit messages, ≤50 char subject. Why over what.
<code>/caveman-review</code>	One-line PR comments: L42: ● bug: user null. Add guard.
<code>/caveman-stats</code>	Real session token usage + lifetime savings + USD. Tweetable line via <code>--share</code> .
<code>/caveman-compress <file></code>	Rewrite memory file (e.g. <code>CLAUDE.md</code>) into caveman-speak. Cuts ~46% input tokens every session. Code/URLs/paths byte-preserved.
<code>caveman-shrink</code>	MCP middleware. Wraps any MCP server, compresses tool descriptions. npm .
<code>cavecrew-*</code>	Caveman subagents (investigator/builder/reviewer). ~60% fewer tokens than vanilla, main context lasts longer.

Statusline badge — Claude Code shows `[CAVEMAN] ↯ 12.4k` (lifetime tokens saved). Updates every `/caveman-stats` run. Set `CAVEMAN_STATUSLINE_SAVINGS=0` to silence.

Auto-activate every session: Claude Code, Codex, Gemini (built-in). Cursor / Windsurf / Cline / Copilot get always-on rule files via `--with-init`. Other agents trigger with `/caveman` per session. Full feature matrix in [INSTALL.md](#).

Benchmarks

Real token counts from the Claude API. Average **65% output reduction** across 10 prompts (range 22-87%).

Task	Normal	Caveman	Saved
Explain React re-render bug	1180	159	87%
Fix auth middleware token expiry	704	121	83%
Set up PostgreSQL connection pool	2347	380	84%
Explain git rebase vs merge	702	292	58%
Refactor callback to async/await	387	301	22%
Architecture: microservices vs monolith	446	310	30%
Review PR for security issues	678	398	41%
Docker multi-stage build	1042	290	72%
Debug PostgreSQL race condition	1200	232	81%
Implement React error boundary	3454	456	87%
Average	1214	294	65%

Raw data and reproduction script: [benchmarks/](#) . Three-arm eval harness (baseline / terse / skill) lives in [evals/](#) — caveman compared against Answer concisely. not against verbose default, so the delta is honest.

caveman-compress receipts (real memory files):

File	Original	Compressed	Saved
claude-md-preferences.md	706	285	59.6%
project-notes.md	1145	535	53.3%
claude-md-project.md	1122	636	43.3%
todo-list.md	627	388	38.1%
mixed-with-code.md	888	560	36.9%
Average	898	481	46%

Caveman only affects output tokens — thinking/reasoning tokens untouched. Caveman no make brain smaller. Caveman make *mouth* smaller. Biggest win is **readability and speed**, cost savings a bonus.

A March 2026 paper "[Brevity Constraints Reverse Performance Hierarchies in Language Models](#)" found that constraining large models to brief responses **improved accuracy by 26 points** on certain benchmarks. Verbose not always better. Sometimes less word = more correct.

How It Work

1. Install drop skill file in agent.
2. Skill tell agent: drop filler, keep substance, use fragments.
3. For Claude Code, hook also write tiny flag file each session — agent see flag, talk caveman from message one. No need say `/caveman` .
4. Stats command read Claude Code session log, count tokens saved, write number to statusline.
5. Caveman-compress sub-skill rewrite memory files (CLAUDE.md, project notes) so each session start with smaller context. Save tokens forever, not just one reply.

Maintainer detail (hook architecture, file ownership, CI sync) live in [CLAUDE.md](#).

Lobster, Meet Rock 🦞 🤖

[OpenClaw](#) the self-host gateway. One box, many agent inside (Claude Code, Codex, Pi, OpenCode), wired to your Slack / Discord / iMessage / Telegram / whatever. Tagline: "*The lobster way*." Lobster strong. Lobster smart. Lobster also talk a lot.

Caveman teach lobster brevity — same canonical installer, scoped to one agent:

```
# macOS / Linux / WSL
curl -fsSL https://raw.githubusercontent.com/JuliusBrussee/caveman/main/install.sh | bash -s

# Windows (PowerShell): no Node? install Node ≥18 first, then
npm -y github:JuliusBrussee/caveman -- --only openclaw
```



Two thing happen, no more:

1. **Skill drop** at `~/.openclaw/workspace/skills/caveman/SKILL.md` — spec-correct frontmatter (`version` , `always: true`), discoverable by `openclaw skills list` . Skill not auto-inject (OpenClaw load skill on demand) — that why we also do step 2.
2. **SOUL.md nudge**. Tiny marker-fenced block appended to `~/.openclaw/workspace/SOUL.md` . OpenClaw inject SOUL.md into *every* turn under "Project Context" (12K-per-file, 60K total — block well under). Lobster terse from message one. No `/caveman` per session. No nag.

~/.openclaw/workspace/

└─ skills/caveman/SKILL.md

└─ SOUL.md

← full ruleset, on-demand load

← <!-- caveman-begin --> ... <!-- caveman-end -->

↑ auto-inject every turn

Custom workspace path? `OPENCLAW_WORKSPACE=/your/path` before the command. Uninstall: same one-liner with `--uninstall` — skill folder gone, SOUL.md block ripped out cleanly, your other workspace content stay untouched. Idempotent re-runs (frontmatter not double-prepended, marker block not duplicated).

Lobster claw still sharp. Lobster mouth now small. Brain still big.

Caveman Ecosystem

Five tools. One philosophy: **agent do more with less**.

Repo	What
caveman (you here)	Output compression — <i>why use many token when few do trick</i>
caveman-code	Whole terminal coding agent — <i>why use many token when whole agent can save</i>
cavemem	Cross-agent memory — <i>why agent forget when agent can remember</i>
cavekit	Spec-driven build loop — <i>why agent guess when agent can know</i>
cavegemma	Gemma 4 31B fine-tuned on caveman pairs — <i>why prompt every turn when weight remember</i>

Compose: cavekit drive build, caveman compress what agent *say*, cavemem compress what agent *remember*, cavegemma bake compression into weight, caveman-code ship it all as one terminal agent. One rock. Two rock. Three rock. Four rock. Five rock. That it.

More Skill From Same Cave

Caveman has siblings. [JuliusBrussee/skills](#) — five skills, one install, works in Claude Code, Cursor, Gemini, Cline, Copilot, 40+ agents:

Skill	What
caveman	This one. Speak less, say more.
grill-me	Agent grill your plan <i>before</i> you build wrong thing. Checks how much you know first — no condescend, no coddle.
interface-kit	Build UI that look good, load fast, work for everyone.

Skill	What
junior-to-senior	Adversarial review pass. Junior output go in, senior output come out.
loop-factory	Spec-driven task loop — inbox → active → archive, review gate between.

```
npx skills@latest add JuliusBrussee/skills
```




One command. Five skill. Cave well stocked.

Links

- [INSTALL.md](#) — full install matrix, all flags, per-agent detail
- [CONTRIBUTING.md](#) — how to send patch
- [CLAUDE.md](#) — maintainer guide (file ownership, hook architecture, CI)
- [docs/](#) — extra guides (Windows install, etc.)
- [Issues](#) — bug, feature, weird behavior

Star This Repo

Caveman save you token, save you money. Star cost zero. Fair trade. ★



Also by Julius Brussee

- [Revu](#) — local-first macOS study app with FRS spaced repetition. [revu.cards](#)

Sponsors

caveman free forever. Sponsor keep rock sharp.

Atlas Cloud

[Atlas Cloud](#) — full-modal AI inference platform, one API.

Want rock here too? [Sponsor caveman](#).

License

MIT — free like mass mammoth on open plain.